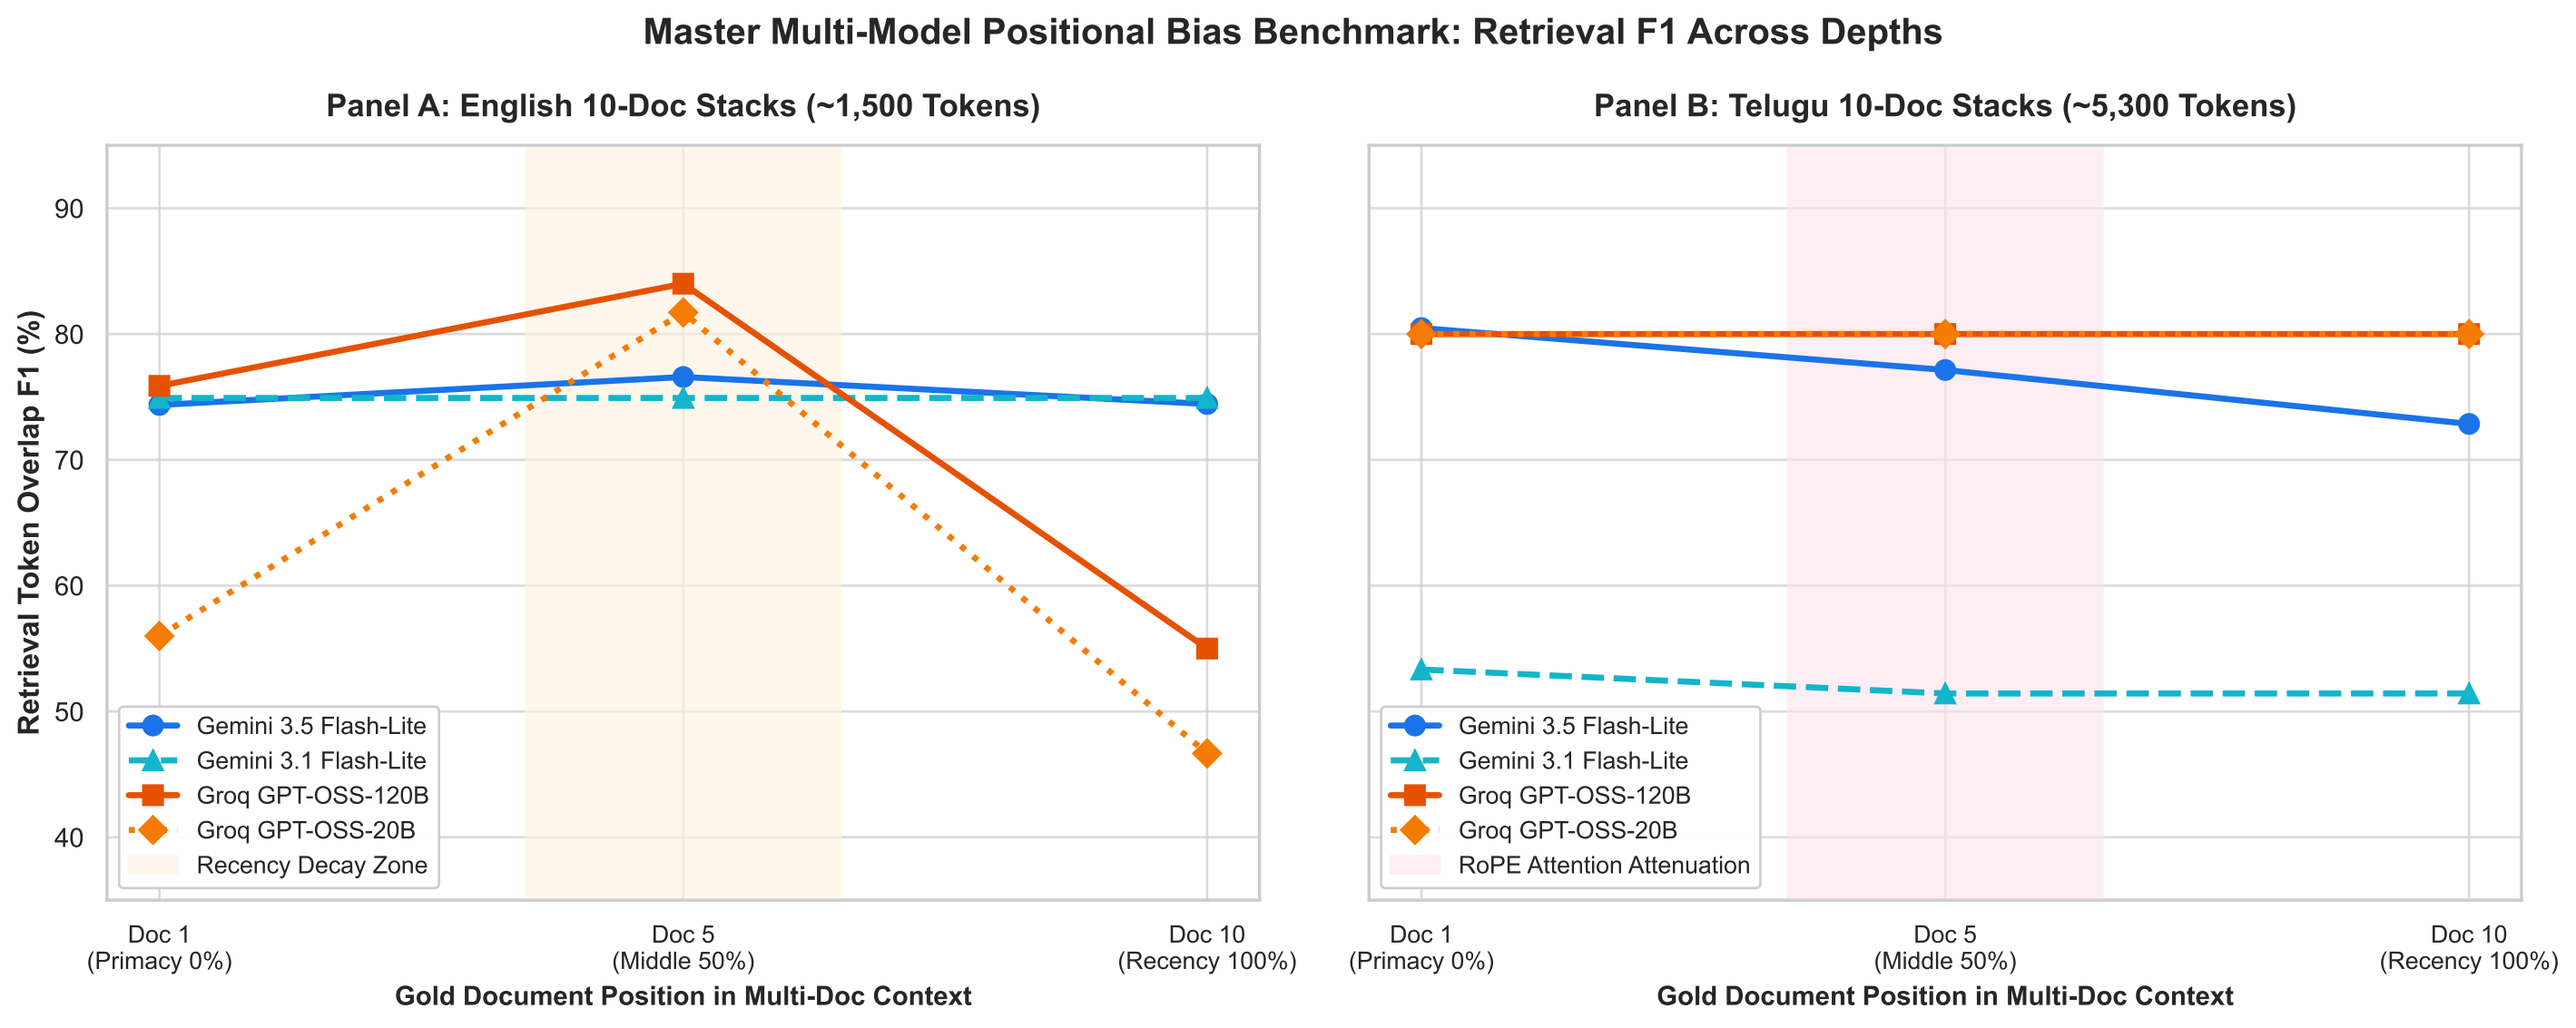
Master Multi-Model Positional Bias Benchmark: Retrieval F1 Across Depths
Panel A: English 10-Doc Stacks (~1,500 Tokens)
Panel B: Telugu 10-Doc Stacks (~5,300 Tokens)
Retrieval Token Overlap F1 (%)
90
80
70
60
50
40
Gemini 3.5 Flash-Lite
Gemini 3.1 Flash-Lite
Groq GPT-OSS-120B
Groq GPT-OSS-20B
Recency Decay Zone
RoPE Attention Attenuation
Doc 1
(Primacy 0%)
Doc 5
(Middle 50%)
Doc 10
(Recency 100%)
Gold Document Position in Multi-Doc Context